

Action Spotting in Soccer Broadcast Videos

CS518: Deep Learning for Computer Vision

Abhinav Sandru · Gaurav Chintakunta · Raviteja Ravella

GitHub: <https://github.com/AbhinavSandru/cs518-action-spotting>

1. Introduction

Broadcast soccer matches are rich sources of tactical and statistical information, but extracting structured event data from raw video requires significant manual effort. **Action spotting** is the task of automatically identifying the precise timestamps of predefined actions in untrimmed long-form video. Unlike action recognition, which classifies a pre-segmented clip, action spotting must locate when events occur within a full 90-minute match without any prior segmentation.

This project addresses action spotting on the SoccerNet-v2 benchmark [1], which defines 17 action classes (goals, fouls, cards, corners, shots, substitutions, etc.) across 500 professional broadcast soccer games. The core challenge is **extreme temporal sparsity**: actions occur at fewer than one frame per thousand per class, making naive classification approaches degenerate — a model can achieve near-zero loss by predicting nothing.

We propose a lightweight **Transformer encoder** trained on precomputed ResNet-152 [2] appearance features, combined with focal loss [3] and inverse-frequency class weighting to handle the severe class imbalance. The full pipeline — preprocessing, training, inference, and evaluation — runs on a single Apple M4 Pro laptop using Metal Performance Shaders (MPS) acceleration.

Final results on the SoccerNet-v2 test set: 38.18% mAP@1s · 52.86% mAP@5s.

2. Related Work / Background / Existing Approaches

2.1 Action Recognition vs. Action Spotting

Standard action recognition (e.g., on Kinetics [4]) classifies pre-trimmed clips. Action spotting is strictly harder: the model must both detect and localize events in untrimmed video, producing (timestamp, class) pairs rather than a single label.

2.2 SoccerNet Benchmark

SoccerNet [1] introduced large-scale soccer action spotting. SoccerNet-v2 [5] extended the dataset to 500 games and 17 classes with tighter annotation quality. Evaluation uses **Average-mAP** — the mean of per-class average precision scores computed at multiple temporal tolerances (1s, 2s, 5s).

2.3 Existing Approaches

NetVLAD++ [6] is the official SoccerNet-v2 baseline. It uses a NetVLAD pooling layer to aggregate temporal context from sliding windows over precomputed features, achieving ~49% mAP@5s.

CALF [7] adds a calibration loss that enforces temporal ordering of predictions, reaching ~42% mAP@5s on SoccerNet-v1.

E2E-Spot [8] trains an end-to-end model directly on raw video frames using a lightweight CNN backbone, achieving 60%+ mAP on SoccerNet-v2 by learning sport-specific features rather than relying on ImageNet-pretrained representations.

Our approach differs from NetVLAD++ by using a Transformer encoder with self-attention instead of learnable pooling, and from E2E-Spot by operating on frozen precomputed features — making it feasible to train on a single laptop. We incorporate focal loss (from object detection [3]) into the action spotting domain, which significantly improves performance on rare classes.

3. Methodology

3.1 Problem Statement

Input: A full soccer broadcast match represented as precomputed ResNet-152 feature vectors $\mathbf{x}_t \in \mathbb{R}^{2048}$ at $t = 1, \dots, T$ frames, where $T \approx 10,800$ per half at 2fps.

Output: A ranked list of predictions $\{(h_i, p_i, c_i, s_i)\}$ where $h_i \in \{1, 2\}$ is the match half, p_i is the timestamp in milliseconds, $c_i \in \{1, \dots, 17\}$ is the action class, and $s_i \in [0, 1]$ is the confidence score.

Per-frame formulation: The model performs binary classification at every frame for each of 17 classes independently. Timestamps are derived post-hoc via non-maximum suppression on the per-frame score arrays.

3.2 Preprocessing

Raw feature files for each game are concatenated across both halves to form a single array of shape $(T, 2048)$. A per-frame binary label array of shape $(T, 17)$ is constructed from the JSON annotations.

To avoid penalizing the model for being one frame off on a 2fps signal, we apply **Gaussian label spreading**: each annotated frame t^* receives a soft target using a Gaussian kernel with $\sigma = 1$ over a ± 2 frame neighborhood:

$$y_{t,c} = \exp\left(-\frac{(t-t^*)^2}{2\sigma^2}\right) \quad \text{for } |t-t^*| \leq 2$$

The full game arrays are then sliced into overlapping 60-frame windows (stride 30) and saved as compressed `.npz` files. This **preprocessing step** reduces training I/O by $8\times$ compared to loading full game files on the fly.

3.3 Model Architecture

The model, **ActionSpotter**, is a standard Transformer encoder operating on sliding windows of frame features.

Step 1 — L2 Normalization: Input features are L2-normalized along the feature dimension, projecting each vector onto the unit hypersphere. This removes scale variance across different games, stadiums, and broadcast cameras.

Step 2 — Linear Projection: A learned linear layer maps from 2048 to $d_{\text{model}} = 256$ dimensions.

Step 3 — Positional Embedding: A learned positional embedding of shape $(1, 1000, 256)$ is added, encoding the temporal position of each frame within the window.

Step 4 — Transformer Encoder: Four encoder layers, each with: - Multi-head self-attention: $n_{\text{head}} = 4$ heads - Feed-forward network: $d_{\text{ff}} = 1024$ - Dropout: $p = 0.3$ - Layer normalization (post-attention)

Step 5 — Classification Head: A linear layer maps each frame’s hidden state to 17 logits.

Output: $(B, T, 17)$ logits — one score per frame per class. Total parameters: ~ 1.4 million.

The key advantage of self-attention over RNNs is that every frame in the 30-second window has direct access to every other frame in $O(1)$ operations, enabling the model to learn long-range temporal patterns (e.g., buildup to a goal).

3.4 Loss Function

With a negative-to-positive ratio of $\$1000:1$ per class, standard binary cross-entropy (BCE) produces degenerate solutions — predicting zero for all frames achieves near-zero loss while detecting nothing.

We combine three strategies:

Focal Loss [3]: For each frame-class pair with predicted probability p_t :

$$\mathcal{L}_{\text{focal}} = -(1 - p_t)^\gamma \log(p_t), \quad \gamma = 2$$

The factor $(1 - p_t)^2$ down-weights easy negatives (where $p_t \approx 0$ and the model is correctly confident), forcing the gradient to concentrate on uncertain, hard positive frames.

Sqrt Inverse-Frequency Class Weights: Per-class positive weight:

$$w_c = \text{clip}\left(\sqrt{\frac{N_c^-}{N_c^+}}, 1, 50\right)$$

where N_c^+ and N_c^- are positive and negative frame counts for class c in the training set. The square-root dampens the raw ratio (\$1000\$)toastablerange(\$8\$–\$45\$). The clip at 50 prevents any single class from dominating the gradient.

Combined loss:

$$\mathcal{L} = \frac{1}{BT} \sum_{b,t} (1 - p_t)^2 \cdot \text{BCE}(\text{logit}_{b,t}, y_{b,t}, w_c)$$

3.5 Training Configuration

Hyperparameter	Value
Optimizer	Adam
Learning rate	1×10^{-4}
LR scheduler	ReduceLROnPlateau (factor 0.5, patience 3)
Early stopping patience	5 epochs
Batch size	32
Gradient clipping	max norm 1.0
Window size	60 frames (30 seconds)
Stride	30 frames
d_{model}	256
Attention heads	4
Encoder layers	4
Dropout	0.3
Focal loss γ	2

Training runs on Apple M4 Pro via PyTorch MPS backend. Convergence typically occurs in 15–20 epochs (~3 hours total).

3.6 Inference

At test time the full game is processed via sliding window inference. For each window position, per-frame probabilities are accumulated and averaged across overlapping windows. **Non-Maximum Suppression (NMS)** is then applied per class: the peak-scoring frame above a threshold is selected as a prediction, and all frames within ± 10 frames are suppressed before repeating.

Per-class thresholds are tuned on the validation set by sweeping $[0.05, 0.50]$ in steps of 0.05 and selecting the threshold maximizing per-class val mAP. Most classes converge to 0.05, reflecting the model’s calibrated but low-magnitude output scores.

4. Experiments

4.1 Ablation Study

We trained incrementally adding each component to measure its individual contribution:

Configuration	mAP@1s	mAP@5s
Baseline: plain BCE, no weighting	~0%	~0%
+ Sqrt inverse-frequency class weights	~24%	~39%

Configuration	mAP@1s	mAP@5s
+ Focal loss ($\gamma = 2$)	~28%	~44%
+ L2 feature normalization	31.94%	~48%
+ Per-class threshold tuning	38.18%	52.86%

The baseline produces zero detections — the model learns to predict nothing. Class weights alone recover meaningful performance. Focal loss adds further improvement by redirecting gradient toward hard examples. L2 normalization is a single-line addition that removes scale variance and gives a consistent gain. Per-class threshold tuning provides the largest single jump.

4.2 Per-Class Results (Test Set)

Class	mAP@1s	mAP@5s
Substitution	0.6951	0.8630
Foul	0.6281	0.8096
Yellow card	0.5572	0.7663
Offside	0.5807	0.6409
Throw-in	0.4852	0.6684
Yellow->red card	0.4757	0.6566
Corner	0.4514	0.6606
Goal	0.4483	0.7390
Goalkeeper saves	0.4422	0.6970
Shots off target	0.4083	0.4636
Shots on target	0.3635	0.4496
Indirect free-kick	0.3262	0.5699
Clearance	0.2850	0.3802
Direct free-kick	0.2825	0.5207
Kick-off	0.0472	0.0745
Red card	0.0143	0.0271
Ball out of play	0.0000	0.0000
Average-mAP	0.3818	0.5286

High performers: Substitution (69.5%) and Foul (62.8%) are visually distinctive — substitutions involve a player in a bib at the sideline with a board held up; fouls involve players clustering and falling. These consistent visual cues are well-captured by ResNet-152 features.

Low performers: Ball out of play (0%) is visually indistinguishable from normal play at 2fps — the ball’s position relative to the line cannot be determined from appearance features alone. Red card (1.4%) is extremely rare (fewer than 1 per game), and Kick-off (4.7%) looks identical to open play.

4.3 Gap Analysis: mAP@1s vs. mAP@5s

The 14-point gap between tight (38.18%) and loose (52.86%) tolerance is largely attributable to the 2fps feature extraction rate. At 2fps, the maximum achievable temporal precision is 0.5 seconds per frame — a structural limitation that affects the 1-second tolerance metric directly. Higher-fps features would close this gap significantly.

5. Conclusion and Future Work

We presented a lightweight Transformer encoder for temporal action spotting on SoccerNet-v2, trained entirely on precomputed ResNet-152 features on a consumer laptop. The key contributions are:

1. Demonstrating that focal loss combined with sqrt inverse-frequency class weighting resolves the degenerate training failure caused by extreme temporal sparsity.
2. Achieving 38.18% mAP@1s and 52.86% mAP@5s — competitive with NetVLAD++ — with only ~1.4M parameters and no end-to-end backbone training.
3. A complete reproducible pipeline including preprocessing, training, inference, threshold tuning, and evaluation.

Future work:

- **End-to-end backbone fine-tuning:** Fine-tuning ResNet-152 on soccer-specific data would produce domain-adapted features and likely close the gap to top methods (60%+).
 - **Multi-modal features:** Adding optical flow or audio would enable detection of visually ambiguous classes like Ball out of play.
 - **Longer temporal context:** Extending beyond 30-second windows using hierarchical attention or full-game BERT-style pretraining could capture match-state context (score, time, fatigue).
 - **Higher frame rate:** Using 25fps features instead of 2fps would remove the structural precision bottleneck at 1-second tolerance.
-

References

- [1] Giancola, S., et al. “SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos.” CVPR Workshops, 2018.
 - [2] He, K., et al. “Deep Residual Learning for Image Recognition.” CVPR, 2016.
 - [3] Lin, T.Y., et al. “Focal Loss for Dense Object Detection.” ICCV, 2017.
 - [4] Kay, W., et al. “The Kinetics Human Action Video Dataset.” arXiv:1705.06950, 2017.
 - [5] Delière, A., et al. “SoccerNet-v2: A Dataset and Benchmarks for Holistic Understanding of Broadcast Soccer Videos.” CVPR Workshops, 2021.
 - [6] Arun, M., et al. “NetVLAD++: Efficient Aggregation of Multi-Frame-Multi-Resolution Features for Action Spotting.” CVPR Workshops, 2021.
 - [7] Cioppa, A., et al. “A Context-Aware Loss Function for Action Spotting in Soccer Videos.” CVPR, 2020.
 - [8] Hong, J., et al. “Spotting Temporally Precise, Fine-Grained Events in Video.” ECCV, 2022.
-

Appendix

A. Team Contribution

This project was completed by Abhinav Sandru, Gaurav Chintakunta, and Raviteja Ravella. All components — dataset setup, preprocessing pipeline, model architecture, training loop, evaluation, threshold tuning, and demo notebook — were implemented collaboratively by the team.

B. Implementation Details

Hardware: Apple MacBook Pro M4 Pro, 24GB unified memory. Training uses PyTorch MPS backend.

Data storage: SoccerNet features (~200GB) stored on external Samsung T5 SSD. Preprocessed `.npz` windows (~70GB) stored on internal SSD for faster training I/O.

Preprocessing: Run once via `python -m src.preprocess`. Generates ~200K train windows, ~65K val/test windows.

Training command: `python main.py --mode train`

Evaluation command: `python main.py --mode evaluate`

Threshold tuning: `python -m src.tune_thresholds`

C. Approaches That Did Not Work

Hard BCE with raw class weights: Using the raw negative/positive ratio ($\sim 1000\times$) as `pos_weight` directly caused gradient explosion — training loss spiked to NaN within 2 epochs. The sqrt scaling was critical.

Global threshold (0.5): Applying a single threshold of 0.5 to all classes produced near-zero predictions since the model's output scores rarely exceed 0.3–0.4. Reducing to 0.05 per-class was necessary.

Large batch sizes with pin_memory: On MPS, `pin_memory=True` causes a warning and slower transfer. Setting `pin_memory=False` and `num_workers=0` gave stable training.

Eager loading (SoccerNetDataset): Loading full game `.npy` files during training caused 18GB+ RAM usage and OOM kills. The preprocessed `.npz` window approach reduced peak RAM to under 4GB.